

SkillDzire - Generative AI with LLMs

1. What does AI primarily aim to achieve?

- A. Increase hardware size
- B. Simulate human intelligence**
- C. Replace databases
- D. Improve internet speed

2. Which of the following is a subset of AI?

- A. Networking
- B. Operating Systems
- C. Compiler Design
- D. Machine Learning**

3. Deep Learning mainly uses:

- A. Linked Lists
- B. Decision Tables
- C. Neural Networks**
- D. Binary Trees

4. NLP stands for:

- A. Neural Language Processing
- B. Natural Language Processing**
- C. Natural Learning Protocol
- D. Network Language Platform

5. Which task belongs to NLP?

- A. Speech Recognition**
- B. Image Compression
- C. CPU Scheduling
- D. Memory Allocation



6. Generative AI is mainly used for:

- A. Data deletion
- B. Content creation**
- C. File encryption
- D. Network routing

7. Which is an example of Generative AI output?

- A. Firewall logs
- B. Cache memory
- C. IP Address
- D. AI-generated text**

8. AI systems learn patterns from:

- A. Random guessing
- B. Data**
- C. Hardware cables
- D. BIOS

9. Which company created the GPT models?

- A. Intel
- B. IBM
- C. OpenAI**
- D. Oracle

10. Which of the following is NOT an AI application?

- A. Chatbots
- B. Recommendation systems
- C. Autonomous vehicles
- D. Manual typewriter**

11. Machine Learning models improve through:

- A. Rebooting
- B. Experience and data**



- C. Cable replacement
- D. Manual wiring

12. Deep Learning performs well with:

- A. Large datasets**
- B. Small data only
- C. No data
- D. Random values

13. Which AI field focuses on human language?

- A. Robotics
- B. NLP**
- C. IoT
- D. Networking

14. Which is an example of supervised learning?


- A. Classification**
- B. Clustering
- C. Dimensionality reduction
- D. Association rules

15. AI agents generally:

- A. Ignore environments
- B. Act autonomously**
- C. Remove RAM
- D. Disable software

16. What is a token in LLMs?

- A. Hardware chip
- B. CPU cycle
- C. Small text unit**
- D. Database key



17. The context window refers to:

- A. Browser size
- B. Number of training GPUs
- C. Amount of text model remembers at once**
- D. RAM slot count

18. Which architecture powers modern LLMs?

- A. CNN only
- B. Transformer**
- C. RNN only
- D. Decision Tree

19. Attention mechanism helps models:

- A. Reduce electricity
- B. Increase monitor brightness
- C. Format hard disks
- D. Focus on relevant words**

20. Embeddings represent:

- A. Images only
- B. Numeric vector representations**
- C. CPU registers
- D. Cache blocks

21. Similar embeddings usually indicate:

- A. Similar meaning**
- B. System crash
- C. Network error
- D. Low RAM

22. Which preprocessing step converts text to lowercase?

- A. Tokenization
- B. Stemming



C. Normalization

D. Parsing

23. Tokenization means:

A. Splitting text into smaller units

B. Encrypting text

C. Compressing files

D. Deleting stop words

24. Stop words are:

A. Important keywords

B. Common words with low meaning contribution

C. Database records

D. Encryption keys

25. Stemming reduces words to:

A. Synonyms

B. Numbers

C. Paragraphs

D. Root forms

26. Lemmatization differs from stemming because it:

A. Uses dictionary meaning

B. Deletes files

C. Increases RAM

D. Changes images

27. Bag of Words represents text using:

A. Images

B. Word frequency

C. Video frames

D. Audio waves



28. TF-IDF helps identify:

- A. Network protocols
- B. Hard disk errors
- C. CPU temperature

D. Important words in documents

29. High TF value means:

A. Word appears frequently in a document

- B. Low storage
- C. Weak internet
- D. Broken token

30. IDF reduces importance of:

- A. Rare words
- B. Common words**
- C. Numbers
- D. Sentences

31. Which is an example of text generation?


- A. Cable routing
- B. Disk partitioning
- C. AI writing an article**
- D. OS booting

32. Prompt engineering means:

- A. Designing hardware prompts
- B. Crafting effective AI instructions**
- C. Formatting drives
- D. Building CPUs

33. Temperature parameter in LLMs controls:

- A. GPU heat
- B. Battery voltage



C. Internet speed

D. Output randomness

34. Higher temperature usually gives:

A. More deterministic responses

B. More creative responses

C. Lower memory

D. Faster internet

35. RAG stands for:

A. Retrieval-Augmented Generation

B. Random AI Generation

C. Retrieval AI Graph

D. Recursive Augmented Grammar

36. Vector databases mainly store:

A. Images only

B. CPU instructions

C. Embeddings

D. BIOS files

37. Similarity search is used in:

A. Kernel compilation

B. BIOS setup

C. Disk formatting

D. RAG systems

38. Chunking in RAG means:

A. Breaking documents into smaller parts

B. Deleting tokens

C. Compressing RAM

D. Reducing CPU speed

39. Which open-source LLM is widely known?

- A. Llama**
- B. Photoshop
- C. MySQL
- D. Excel

40. Fine-tuning modifies:

- A. Hardware drivers
- B. Model parameters using training data**
- C. Browser tabs
- D. Internet cables

41. Prompting is often cheaper than:

- A. Networking
- B. Parsing
- C. Compression
- D. Fine-tuning**

42. AI evaluation measures:

- A. Keyboard size
- B. Model performance**
- C. CPU wiring
- D. Battery color

43. Hallucination in LLMs refers to:

- A. GPU overheating
- B. Incorrect generated information**
- C. Slow internet
- D. Broken APIs

44. Streamlit is commonly used for:

- A. AI web app deployment**
- B. PCB manufacturing



C. File encryption

D. Game engines

45. Gradio helps developers:

A. Create quick AI interfaces

B. Design processors

C. Build operating systems

D. Repair hardware

46. Which component retrieves relevant documents in RAG?

A. Retriever

B. BIOS

C. Router

D. Compiler

47. Embedding models convert text into:

A. Videos

B. Vectors

C. Tables

D. Circuits

48. Which task is commonly solved using NLP?

A. Machine lubrication

B. Sentiment analysis

C. RAM soldering

D. Signal boosting

49. AI chatbots rely heavily on:

A. NLP

B. Hydraulic systems

C. BIOS

D. Mechanical relays

50. Mini AI projects help learners:

- A. Avoid coding
- B. Apply concepts practically**
- C. Disable training
- D. Remove datasets

51. Artificial Intelligence is the broader concept of:

- A. Human labor
- B. Smart machines**
- C. Mechanical repair
- D. Hardware assembly

52. Which Machine Learning approach learns using trial and error interactions?

- A. Supervised Learning
- B. Unsupervised Learning
- C. Reinforcement Learning**
- D. Batch Learning

53. Which layer in a neural network produces the final prediction?

- A. Hidden layer
- B. Convolution layer
- C. Output layer**
- D. Embedding layer

54. Which learning type uses labeled data?

- A. Supervised learning**
- B. Unsupervised learning
- C. Reinforcement learning
- D. Random learning

55. Which learning type finds hidden patterns?

- A. Supervised learning
- B. Unsupervised learning**
- C. Static learning
- D. Manual learning

56. Reinforcement learning uses:

A. Rewards and penalties

B. Word counts

C. IP addresses

D. Database indexing

57. Which hardware component is most commonly used to accelerate Deep Learning training?

A. Sound card

B. GPU

C. Network switch

D. Optical drive

58. Which algorithm is commonly used in ML?

A. Linear Regression

B. HTTP

C. DNS

D. FTP

59. CNNs are especially useful for:

A. Database backups

B. Image processing

C. Email routing

D. BIOS management

60. RNNs are suitable for:

A. USB transfer

B. Power supply

C. LAN wiring

D. Sequential data

61. Which of the following is NOT Deep Learning?

A. Transformer

B. CNN

C. Decision Tree

D. RNN

62. Training data is used to:

A. Build hardware

B. Teach models patterns

C. Increase RAM physically

D. Format drives

63. Overfitting occurs when a model:

A. Shuts down

B. Stops training early

C. Deletes records

D. Learns training data too specifically

64. Underfitting means:

A. Model performs poorly on data

B. GPU overload

C. Excess RAM usage

D. Internet failure

65. Which metric is used for classification evaluation?

A. Bandwidth

B. Voltage

C. Accuracy

D. Clock speed

66. Precision measures:

A. Total memory

B. Correct positive predictions

C. CPU cycles

D. Training speed only

67. Recall measures:

A. Disk usage

B. Captured actual positives

C. Battery health

D. GPU size

68. F1-score combines:

A. Precision and Recall

B. Voltage and Current

C. CPU and RAM

D. Token and chunk

69. Gradient descent is used for:

A. Encryption

B. Compression

C. Optimization

D. Parsing

70. Epoch in ML means:

A. Internet refresh

B. CPU restart

C. GPU shutdown

D. One full training cycle through data

71. Batch size refers to:

A. Number of training examples processed together

B. RAM slots

C. Hard disks

D. Browser windows

72. Learning rate controls:

A. Speed of parameter updates

B. Internet bandwidth



- C. GPU temperature
- D. Hard disk RPM

73. Activation functions introduce:

- A. File deletion
- B. Cabling
- C. Non-linearity**
- D. Compression only

74. ReLU stands for:

- A. Random Linear Unit
- B. Rectified Linear Unit**
- C. Reduced Logic Utility
- D. Recursive Learning Unit

75. Sigmoid activation outputs values between:

- A. -10 and 10
- B. 0 and 1**
- C. 1 and 100
- D. -1 and -10

76. Softmax is commonly used in:

- A. Disk formatting
- B. CPU testing
- C. Multi-class classification**
- D. Image compression only

77. Backpropagation helps:

- A. Update neural network weights**
- B. Increase voltage
- C. Connect routers
- D. Format memory cards

78. GPUs accelerate:

- A. Keyboard typing
- B. Deep Learning computations**
- C. Mouse clicks
- D. BIOS loading

79. Transfer learning uses:

- A. BIOS settings
- B. Network switches
- C. Random cables
- D. Pretrained models**

80. Feature engineering means:

- A. Creating useful input variables**
- B. Building CPUs
- C. Repairing drives
- D. Upgrading routers

81. Which library is popular for Deep Learning?

- A. VLC
- B. TensorFlow**
- C. Chrome
- D. WinRAR

82. PyTorch is widely used for:

- A. Audio routing
- B. Video editing
- C. File compression
- D. Deep Learning**

83. Bias in ML refers to:

- A. Fan speed
- B. USB port

C. Systematic error

D. Cache memory

84. Variance relates to:

A. Sensitivity to training data changes

B. Voltage levels

C. RAM type

D. Browser cache

85. Cross-validation improves:

A. CPU cooling

B. Model reliability

C. Hard disk rotation

D. Internet access

86. Confusion matrix is used for:

A. Classification evaluation

B. Database design

C. GPU cooling

D. Encryption keys

87. Which learning algorithm is unsupervised?

A. Decision Tree

B. Logistic Regression

C. K-Means

D. SVM

88. Decision Trees are mainly used for:

A. PCB design

B. Hardware repair

C. Classification and regression

D. Network routing

89. Random Forest is:

A. Ensemble method

B. Virus scanner

C. Router firmware

D. File manager

90. SVM stands for:

A. Signal Voltage Module

B. Storage Virtual Memory

C. Support Vector Machine

D. Smart Visual Monitor

91. Clustering groups:

A. CPUs

B. Similar data points

C. Routers

D. Browsers

92. Dimensionality reduction helps:

A. Reduce feature count

B. Increase monitor size

C. Add RAM slots

D. Expand storage physically

93. PCA stands for:

A. Program Control Adapter

B. Primary CPU Access

C. Parallel Cache Array

D. Principal Component Analysis

94. Neural networks contain:

A. Network cables only

B. Water pipes



C. Layers of neurons

D. Mechanical gears

95. Hidden layers are found in:

A. Neural networks

B. Printers

C. Routers

D. Keyboards

96. Which AI field enables self-driving cars?

A. BIOS engineering

B. File compression

C. Spreadsheet analysis

D. Computer Vision

97. Data imbalance can affect:

A. Screen brightness

B. Model performance

C. CPU sockets

D. USB ports

98. Hyperparameters are:

A. File extensions

B. GPU fans

C. Settings chosen before training

D. Browser cookies

99. AI ethics mainly concerns:

A. Responsible AI usage

B. Hard disk speed

C. Monitor size

D. RAM voltage



100. Explainable AI helps users:

- A. Encrypt databases
- B. Increase internet speed
- C. Reduce hardware cost only

D. Understand model decisions

101. Which model type is mainly used in Generative AI?

- A. Switching circuits
- B. Sorting algorithms
- C. Database schemas

D. Generative models

102. ChatGPT is primarily a:

- A. Spreadsheet tool
- B. Large Language Model application**
- C. Hardware compiler
- D. Network router

103. LLM stands for:

- A. Large Logic Machine
- B. Linear Learning Mechanism
- C. Large Language Model**
- D. Language Linking Module

104. Autocomplete text generation predicts:

- A. Previous token
- B. Next token**
- C. Random image
- D. Hardware address

105. Which company developed BERT?

- A. Microsoft
- B. Google**



- C. Intel
- D. IBM

106. BERT is mainly designed for:

- A. Disk encryption
- B. Browser rendering
- C. Bidirectional understanding**
- D. CPU scheduling

107. GPT models are mainly:

- A. Encoder-only
- B. CNN-only
- C. Decoder-based**
- D. Rule-only

108. Which preprocessing step removes punctuation?

- A. Routing
- B. Cleaning**
- C. Encoding
- D. Scheduling

109. Named Entity Recognition identifies:

- A. Hardware names
- B. IP addresses only
- C. Important entities in text**
- D. RAM modules

110. POS tagging refers to:

- A. Point of Storage
- B. Process Output Scaling
- C. Part-of-Speech tagging**
- D. Packet Operating System

111. Which is an example of sentiment analysis?

- A. Formatting documents
- B. Calculating RAM usage
- C. Building circuits

D. Detecting positive or negative reviews

112. Translation systems are applications of:

A. NLP

- B. Networking
- C. Cybersecurity
- D. BIOS

113. Which technique converts text into vectors?

- A. Compression
- B. Embedding**
- C. Parsing
- D. Routing

114. Cosine similarity measures:

- A. CPU voltage
- B. Hard disk speed
- C. Text similarity between vectors**
- D. Screen resolution

115. The attention mechanism was introduced prominently in:

- A. Relational databases
- B. Operating systems
- C. Transformers**
- D. Routers

116. Self-attention helps a model:

- A. Increase internet speed
- B. Delete duplicate files

C. Optimize storage

D. Focus on relationships between words

117. Transformer models process tokens:

A. Sequentially only

B. In parallel

C. Without memory

D. Using BIOS

118. Which is an advantage of transformers over RNNs?

A. Smaller keyboards

B. Better parallelization

C. Less storage always

D. No training needed

119. Context length in LLMs affects:

A. CPU temperature

B. Monitor size

C. Fan speed

D. Amount of information remembered

120. Which tokenization method splits words into subwords?

A. Rasterization

B. Byte Pair Encoding

C. Encryption

D. Clustering

121. Temperature = 0 in LLMs usually gives:

A. Random outputs only

B. GPU shutdown

C. Slower internet

D. More deterministic output

122. Top-k sampling limits:

A. RAM size

B. Candidate token choices

C. GPU voltage

D. Network ports

123. Prompt templates help:

A. Repair storage devices

B. Increase monitor brightness

C. Standardize AI interactions

D. Replace databases

124. Few-shot prompting provides:

A. BIOS drivers

B. GPU cooling

C. Example inputs and outputs

D. Disk partitions

125. Zero-shot prompting means:

A. Zero RAM usage

B. No tokens used

C. No GPU required

D. No examples are provided

126. Chain-of-thought prompting encourages:

A. Disk formatting

B. Hardware debugging

C. File compression

D. Step-by-step reasoning

127. Hallucination risk in LLMs can be reduced using:

A. RAG

B. Random guessing



- C. File deletion
- D. Lower RAM

128. API stands for:

- A. Automated Processing Input
- B. Applied Program Installation
- C. Application Programming Interface**
- D. Advanced Port Integration

129. OpenAI APIs allow developers to:

- A. Repair hardware
- B. Install operating systems
- C. Manufacture CPUs
- D. Access AI models programmatically**

130. Which parameter limits output length?

- A. Learning rate
- B. Max tokens**
- C. Batch size
- D. Epoch count

131. Prompt injection is mainly a:

- A. Hardware upgrade
- B. Security concern**
- C. Compression technique
- D. Browser extension

132. Token limits are important because:

- A. GPUs stop working
- B. Routers disconnect
- C. Screens freeze
- D. Models have context constraints**

133. Fine-tuning customizes:

- A. Hard disks physically
- B. Keyboard layouts
- C. Internet routers

D. Model behavior using specific data

134. Instruction-tuned models are trained to:

- A. Increase storage space

B. Follow user instructions better

- C. Build operating systems
- D. Encrypt databases

135. RLHF stands for:

- A. Random Learning Hybrid Function
- B. Recursive Logic Hardware Framework
- C. Real-time Linear Hash Function

D. Reinforcement Learning from Human Feedback

136. Which task commonly uses embeddings?

- A. BIOS reset
- B. Semantic search**
- C. Disk formatting
- D. Fan calibration

137. Semantic search focuses on:

- A. File extension names
- B. Folder colors
- C. Hardware serial numbers

D. Meaning-based retrieval

138. Vector dimensions represent:

- A. GPU fans
- B. Internet channels

C. Features in embedding space

D. USB slots

139. Which database is commonly used for vectors?

A. Excel

B. Pinecone

C. Oracle Forms

D. Notepad

140. FAISS is mainly used for:

A. CPU manufacturing

B. Image printing

C. Cable routing

D. Similarity search

141. Chunk overlap in RAG helps:

A. Reduce voltage

B. Increase screen brightness

C. Compress images only

D. Preserve context continuity

142. Retriever models in RAG fetch:

A. Power cables

B. Relevant documents

C. GPU drivers

D. BIOS files

143. Generator models in RAG:

A. Format hard drives

B. Compress vectors

C. Create routers

D. Produce final responses

144. OCR is useful for:

- A. Increasing bandwidth
- B. Cooling processors
- C. Installing RAM

D. Extracting text from images

145. PDF parsing is important in:

- A. Power supply testing
- B. Motherboard assembly
- C. Document processing pipelines**
- D. Router firmware updates

146. Which open-source framework supports LLM apps?

- A. Photoshop
- B. LangChain**
- C. Excel VBA
- D. AutoCAD

147. AI agents can:

- A. Replace electricity
- B. Build GPUs physically
- C. Increase monitor resolution

D. Perform autonomous multi-step tasks

148. Tool calling enables LLMs to:

- A. Repair hardware
- B. Create processors
- C. Use external tools and APIs**
- D. Encrypt SSDs

149. Streamlit applications are typically written in:

- A. Python**
- B. Java Assembly

- C. HTML only
- D. BIOS scripts

150. Gradio is especially useful for:

- A. Database normalization
- B. Hard disk partitioning
- C. PCB manufacturing

D. Rapid AI demo interfaces

151. Which stage in a RAG pipeline converts raw documents into searchable vector representations?

- A. Token decoding

B. Embedding generation

- C. Output rendering
- D. Response ranking

152. Which similarity metric is commonly used with embeddings?

- A. Packet rate

B. Cosine similarity

- C. Voltage drop
- D. Pixel ratio

153. Dense embeddings capture:

- A. CPU frequency
- B. File size only

C. Semantic meaning

- D. Hardware voltage

154. Sparse vectors are commonly associated with:

- A. GPUs

B. TF-IDF

- C. BIOS
- D. Transformers only

155. Which technique improves retrieval quality in RAG?

- A. Deleting embeddings
- B. Reducing RAM
- C. Lowering bandwidth

D. Better chunking

156. Metadata in vector databases helps:

- A. Increase CPU speed
- B. Format SSDs
- C. Filter search results**
- D. Reduce screen size

157. Hybrid search combines:

- A. CPU and RAM
- B. BIOS and drivers
- C. Keyword and vector search**
- D. Images and videos only

158. Which stage converts documents into embeddings?


- A. BIOS loading
- B. Indexing pipeline**
- C. Disk partitioning
- D. File compression

159. Retrieval latency refers to:

- A. GPU temperature
- B. Internet color depth
- C. Time taken to fetch results**
- D. CPU voltage

160. Which file format is commonly processed in RAG systems?

- A. EXE only
- B. DLL only



C. ISO only

D. PDF

161. Chunk size selection affects:

A. Keyboard brightness

B. Retrieval accuracy

C. USB transfer voltage

D. Power supply capacity

162. Embedding normalization improves:

A. GPU fan speed

B. Disk rotation

C. Similarity comparison consistency

D. Browser rendering

163. Which framework supports orchestration of LLM workflows?

A. Photoshop

B. LangChain

C. MySQL Workbench

D. Blender

164. LlamaIndex is mainly designed for:

A. Circuit simulation

B. Browser development

C. Data integration with LLMs

D. Video editing

165. Which model type generates embeddings?

A. Compiler model

B. Embedding model

C. Rendering model

D. BIOS model

166. Open-source LLMs allow:

- A. GPU shutdown
- B. BIOS replacement
- C. File encryption only
- D. Community customization**

167. Quantization mainly helps:

- A. Increase monitor size
- B. Improve keyboard layout
- C. Expand RAM physically
- D. Reduce model size**

168. 4-bit quantization is mainly used for:

- A. Image printing
- B. Hard disk repair
- C. CPU overclocking
- D. Efficient inference**

169. Inference refers to:

- A. Data deletion
- B. Disk cloning
- C. Model prediction process**
- D. Router rebooting

170. GPU memory is important for:

- A. Keyboard input
- B. File naming
- C. Browser themes
- D. Running large AI models**

171. Which deployment tool quickly creates AI dashboards?

- A. BIOS updater
- B. Streamlit**



- C. WinZip
- D. VLC

172. Gradio interfaces are commonly used for:

- A. SSD formatting
- B. Hardware diagnostics

C. AI model demos

- D. Network mapping

173. Docker is useful for:

- A. Increasing monitor brightness
- B. BIOS installation
- C. File compression only

D. Containerized deployment

174. REST APIs commonly use:

- A. CPU registers
- B. HTTP requests**
- C. BIOS interrupts
- D. GPU shaders

175. Which HTTP method usually retrieves data?

- A. GET**
- B. DELETE
- C. PATCH
- D. TRACE

176. JSON is commonly used for:

- A. Hard disk formatting
- B. Voltage control
- C. API data exchange**
- D. GPU rendering

177. Which cloud platform supports AI deployment?

- A. AWS**
- B. Paint
- C. Notepad
- D. Calculator

178. Latency in AI applications affects:

- A. Monitor resolution
- B. Response speed**
- C. File size only
- D. Printer ink usage

179. Throughput measures:

- A. Voltage level
- B. GPU fan rotation
- C. Number of processed requests**
- D. Screen brightness

180. Which metric measures generated text quality?

- A. FPS
- B. BLEU score**
- C. GHz
- D. DPI

181. ROUGE metric is often used for:

- A. Keyboard testing
- B. Power monitoring
- C. Text summarization evaluation**
- D. SSD optimization

182. Hallucination evaluation checks:

- A. Hard disk errors
- B. Router voltage

C. Factual correctness

D. GPU rendering speed

183. Human evaluation in LLMs assesses:

A. CPU fan efficiency

B. SSD speed only

C. Response usefulness and quality

D. Browser themes

184. AI safety focuses on:

A. Screen protection

B. Electrical grounding

C. Responsible model behavior

D. Hardware insulation

185. Bias in generative AI may cause:

A. Faster GPUs

B. Better SSD cooling

C. Unfair outputs

D. Increased voltage

186. Prompt leakage is a:

A. Hardware upgrade

B. File backup method

C. Security concern

D. Browser plugin


187. Jailbreaking an LLM attempts to:

A. Repair operating systems

B. Increase storage

C. Improve RAM speed

D. Bypass safety restrictions



188. Multi-modal AI models can process:

- A. Voltage only
- B. File names only
- C. Text and images**
- D. CPU instructions only

189. Speech-to-text systems convert:

- A. Images into videos
- B. RAM into storage
- C. Audio into text**
- D. HTML into CSS

190. Text-to-speech systems generate:

- A. Hard disk sectors
- B. Browser cookies
- C. GPU firmware
- D. Spoken audio from text**

191. Which AI technique powers recommendation systems?

- A. Voltage regulation
- B. Machine Learning**
- C. PCB soldering
- D. BIOS mapping

192. AI copilots are designed to:

- A. Replace electricity
- B. Format drives
- C. Manufacture CPUs
- D. Assist users interactively**

193. Autonomous AI agents often use:

- A. SSD partitioning
- B. BIOS flashing

C. Planning and tool usage

D. Monitor calibration

194. Memory systems in AI agents help:

A. Increase internet speed

B. Reduce CPU heat

C. Store conversation context

D. Format databases

195. Which protocol is commonly used for AI APIs?

A. HTTPS

B. VGA

C. SATA

D. HDMI

196. Token cost in commercial APIs usually depends on:

A. Screen size

B. CPU brand

C. Number of processed tokens

D. Mouse sensitivity

197. Edge AI refers to:

A. Internet browser edges

B. GPU borders

C. Screen margins

D. Running AI on local devices

198. Which hardware is optimized for AI acceleration?

A. GPU

B. Keyboard

C. Printer

D. Webcam

199. AI benchmarking compares:

- A. Browser wallpapers
- B. Fan colors
- C. Model performance**
- D. Power sockets

200. A mini AI project mainly helps students:

- A. Avoid coding practice
- B. Reduce storage usage only
- C. Replace operating systems
- D. Gain practical implementation skills**

201. System prompts primarily define:

- A. GPU voltage
- B. SSD partitions
- C. Network topology
- D. Model behavior guidelines**

202. User prompts are:

- A. Hardware drivers
- B. Instructions given by users**
- C. BIOS files
- D. Cache blocks

203. Prompt chaining means:

- A. Chaining hard disks
- B. Linking monitors
- C. Connecting multiple prompts sequentially**
- D. Wiring GPUs

204. ReAct prompting combines:

- A. RAM and CPU
- B. Reasoning and acting**



C. Voltage and current

D. Files and folders

205. AI agents typically operate using:

A. Goals and actions

B. BIOS updates only

C. Random clicking

D. Hardware interrupts

206. Multi-agent systems involve:

A. Multiple keyboards

B. Multiple hard disks

C. Multiple AI agents collaborating

D. Multiple GPUs only

207. Agent memory improves:

A. GPU cooling

B. Context retention

C. SSD formatting

D. Screen rendering

208. Planning modules in AI agents help:

A. Increase voltage

B. Repair circuits

C. Expand RAM slots

D. Determine task sequences

209. Tool integration allows agents to:

A. Build processors

B. Repair monitors

C. Access external services

D. Compress SSDs

210. Which component stores long-term AI memory?

- A. Vector store**
- B. GPU fan
- C. Power supply
- D. BIOS chip

211. Autonomous agents can:

- A. Repair motherboards
- B. Replace electricity
- C. Perform tasks without continuous human control**
- D. Increase monitor resolution

212. AI orchestration frameworks help:

- A. Assemble CPUs
- B. Design routers
- C. Coordinate workflows**
- D. Paint interfaces

213. Which prompt style asks the model to think step-by-step?

- A. Binary prompting
- B. Chain-of-thought**
- C. Sparse prompting
- D. Indexed prompting

214. Structured prompting improves:

- A. GPU speed
- B. SSD lifespan
- C. Response consistency**
- D. Keyboard response

215. Which AI deployment style hosts models on cloud servers?

- A. Edge deployment
- B. BIOS deployment

C. Cloud deployment

D. Router deployment

216. Local LLM deployment offers:

A. Better privacy control

B. Larger monitors

C. Faster internet

D. Increased RAM automatically

217. API rate limits restrict:

A. GPU voltage

B. SSD sectors

C. Number of requests allowed

D. Mouse sensitivity

218. Which open-source model family is popular for local deployment?

A. Oracle DB

B. Llama

C. ChromeOS

D. MATLAB

219. Fine-tuning requires:

A. Extra keyboards

B. Larger monitors

C. Additional training data

D. BIOS drivers only

220. LoRA fine-tuning is popular because it:

A. Increases internet speed

B. Reduces training requirements

C. Improves printer quality

D. Repairs SSDs

221. Parameter-efficient tuning aims to:

- A. Reduce screen size
- B. Increase GPU count only
- C. Update fewer model parameters**
- D. Remove APIs

222. Benchmark datasets are used for:

- A. Hardware testing only
- B. Model evaluation**
- C. BIOS flashing
- D. Cable management

223. Which issue affects factual reliability in LLMs?

- A. SSD fragmentation
- B. GPU overheating
- C. Power outage
- D. Hallucination**

224. AI governance mainly addresses:

- A. Keyboard layouts
- B. USB drivers
- C. Policies and responsible use**
- D. Screen calibration

225. Guardrails in AI systems help:

- A. Increase GPU memory
- B. Repair databases
- C. Enforce safety constraints**
- D. Compress files

226. Which format is commonly used for prompt templates?

- A. Structured text placeholders**
- B. Binary machine code

- C. BIOS assembly
- D. GPU shaders

227. Function calling in LLMs enables:

- A. CPU manufacturing
- B. SSD formatting
- C. Structured tool execution**
- D. Keyboard remapping

228. AI copilots are widely used for:

- A. Hardware soldering
- B. Voltage regulation
- C. Network cabling
- D. Productivity assistance**

229. Which platform is widely used for sharing ML models?

- A. Paint
- B. Hugging Face**
- C. BIOS Center
- D. VLC

230. Hugging Face Transformers provides:

- A. Electrical transformers
- B. GPU firmware
- C. Pretrained transformer models**
- D. SSD drivers

231. Which Python library is commonly used for NLP?

- A. NumLock
- B. spaCy**
- C. BIOSLib
- D. GPUBoost

232. NLTK is mainly used for:

A. NLP research and preprocessing

B. Internet routing

C. PCB manufacturing

D. GPU calibration

233. Which AI task summarizes long documents?

A. Packet switching

B. Disk cloning

C. Voltage balancing

D. Text summarization

234. Extractive summarization selects:

A. Random pixels

B. GPU drivers

C. Existing sentences from text

D. Network packets

235. Abstractive summarization:

A. Generates new summary wording

B. Repairs RAM

C. Formats SSDs

D. Increases bandwidth

236. Which NLP task predicts categories?

A. Text classification

B. PCB etching

C. Voltage conversion

D. BIOS loading

237. Spam detection is an example of:

A. GPU rendering

B. Hard disk partitioning



C. Hardware virtualization

D. Text classification

238. Which architecture introduced positional encoding?

A. Transformer

B. Decision Tree

C. CNN

D. K-Means

239. Positional encoding helps models understand:

A. GPU voltage

B. SSD speed

C. Token order

D. File ownership

240. Encoder-decoder transformers are commonly used for:

A. Translation tasks

B. Browser rendering

C. PCB testing

D. Hardware cooling

241. Which training objective predicts masked words?

A. Clustering

B. Regression

C. Masked Language Modeling

D. Routing

242. Which task predicts the next token?

A. Voltage testing

B. BIOS programming

C. Causal Language Modeling

D. Image rendering

243. Beam search is used in:

- A. GPU cooling
- B. Sequence generation**
- C. SSD repair
- D. Circuit routing

244. Greedy decoding selects:

- A. Random token always
- B. Lowest voltage path
- C. Fastest GPU
- D. Highest probability token each step**

245. Sampling-based decoding increases:

- A. CPU frequency
- B. Disk space
- C. Output diversity**
- D. Router range

246. AI observability helps monitor:

- A. Keyboard lighting
- B. Model behavior in production**
- C. Fan rotation
- D. Screen borders

247. Drift detection identifies:

- A. SSD fragmentation
- B. GPU overheating
- C. Changes in incoming data patterns**
- D. BIOS corruption

248. Which concept measures ethical fairness in AI?

- A. Voltage regulation
- B. Bias mitigation**

C. Packet filtering

D. Disk cleanup

249. Explainability tools help developers:

A. Interpret model decisions

B. Upgrade routers

C. Repair processors

D. Compress operating systems

250. Capstone AI projects are valuable because they:

A. Eliminate coding completely

B. Replace datasets

C. Demonstrate practical end-to-end skills

D. Remove deployment needs



Let's Learn and Grow